

Computer Vision II - Homework Assignment 4

Stefan Roth, Xiang Chen, Jochen Gast, Faraz Saeedan, Nikita Araslanov
Visual Inference Lab, TU Darmstadt

July 8, 2019

This homework is due on August 5th, 2019 at 11:00 am.

Please read the instructions carefully!

General remarks

Your grade does not only depend on the correctness of your answer but also on clear presentation of your results and good writing style. It is your responsibility to find a way to *explain clearly how* you solve the problems. Note that we will assess your complete solution and not exclusively the results you present to us. If you get stuck, try to explain why and describe the problems you encounter – you can get partial credit even if you have not completed the task. Hence, please hand in enough information so that we can understand what you have done, what you have tried, and how your final solution works.

Every group has to submit its own original solution. We encourage interaction about class-related topics both within and outside of class. However, you are not allowed to share solutions with your classmates, and *everything you hand in must be your own work*. Also, you are not allowed to just copy material from the web. You are required to *acknowledge any source of information you use to solve the homework* (*i.e.* books other than the course books, papers, websites, etc). Acknowledgments will *not* affect your grade. Not acknowledging a source you rely on is a clear violation of academic ethics. Note that both the university and the department are very serious about plagiarism. For more details, see the department guidelines about plagiarism at https://www.informatik.tu-darmstadt.de/studium_fb20/im_studium/studienbuero/plagiarismus/index.de.jsp and <http://plagiarism.org>.

Programming exercises

For the programming exercises you will be asked to hand in Python code. Please make sure that your code complies with **Python 3.x / PyTorch 1.1**. In order for us to be able to grade the programming assignments properly, stick to the function names and type annotations that we provide in the assignments. Additionally, comment your code in sufficient detail so that it will be easy to understand for us what each part of your code does. Sufficient detail does not mean that you should comment every line of code (that defeats the purpose), nor does it mean that you should comment 20 lines of code using only a single sentence. Your Python code should display your results so that we can judge if your code works from the results alone. Of course, we will still look at the code. If your code displays results in multiple stages, please insert appropriate `sleep` commands between the stages so that we can step through the code. Group plots that semantically belong together in a single figure using subplots and don't forget to put proper titles and other annotations on the plots. Please be sure to name each file according to the naming scheme included with each problem. This also makes it easier for us to grade your submission. And finally, please make sure that you included your name and email in the code.

Files you need

All the data you will need for the problems will be made available in Moodle.

What to hand in

Your hand-in should contain a PDF file (a plain text file is ok, too) with any textual answers that may be required. You must not include images of your results; your code should display these instead. For the programming parts, please hand in all documented `.py` scripts and functions that your solution requires. Make sure your code actually works and that all your results are displayed properly!

Handing in

Please upload your solution files as a single `.zip` or `.tar.gz` file to the corresponding Moodle area at <https://moodle.tu-darmstadt.de/course/view.php?id=15293>. **Please note that we will not accept file formats other than the ones specified!** Your archive should include your write-up (`.pdf` or `.txt`) as well as your code (`.py` scripts). If *and only if* you have problems with your upload, you may send it to `cv2staff@visinf.tu-darmstadt.de`

Late Handins

We will accept late hand-ins, but we will deduct 20% of the total reachable points for every day that you are late. Note that even 15 minutes late will be counted as being one day late! After the exercise has been discussed in class, you can no longer hand in. If you are not able to make the deadline, *e.g.* due to medical reasons, you need to contact us *before* the deadline. We might waive the late penalty in such a case.

Code Interviews

After your submission, we may invite you to give a code interview. In the interview you need to be able to explain your written solution as well as your submitted code to us.

Python Environment

Please follow the instructions in `Readme.txt` to set up your environment.

Problem 1 – Iterative graph cuts for image segmentation

27 points

In this problem we apply graph cuts under a contrast-sensitive Potts model for image segmentation, and extend it for multi-label segmentation via α -expansion similar to the stereo algorithm of the last assignment. The overall approach is based on a simplified version of the GrabCut¹ algorithm for binary foreground-background segmentation.

Binary, interactive GrabCut. Here, we briefly review the original version: GrabCut, an iterative graph cuts algorithm, expects color images as input as well as a bounding box annotation drawn around the object of interest. The algorithm then proceeds to refine a binary labeling for each pixel either classifying it as foreground or background. The unary potentials of the underlying graph cuts algorithm are essentially obtained from learning color statistics of foreground and background pixels, respectively. The pairwise potentials are based on a contrast-sensitive Potts model encouraging labels to respect discontinuities in the color image. Pixels outside the bounding box are fixed to be in the background. With I , Θ and S denoting the image, the parameters of the color distribution and the segmentation, respectively, the segmentation energy for pixels inside the bounding box is given by

$$E(S | I, \Theta) = E_d(S | I, \Theta) + \lambda E_s(S | I), \quad (1)$$

where E_d and E_s are the data (unary) term and the smoothness (pairwise) term, respectively, and λ is a trade-off parameter. The data term

$$E_d(S | I, \Theta) = \sum_{k=1}^K -\log \text{GMM}(I_k | \theta_{S_k}) \quad (2)$$

measures how well each pixel fits to the color distribution of foreground or background pixels (depending on the segmentation). Color distributions are represented as Gaussian Mixture Models with K components. The spatial smoothness term favors segmentations that align with the image boundaries by using contrast sensitive Potts potentials on a 4-connected MRF:

$$E_s(S | I) = \sum_{(i,j) \in \mathcal{N}} \exp(-\beta \|I_i - I_j\|^2) \cdot \delta_{S_i \neq S_j}. \quad (3)$$

The iterative algorithm then proceeds by initializing the color distributions from user annotation, *e.g.* a bounding box, and then iterating between minimizing the segmentation energy and re-fitting the color distributions.

Multi-label, unsupervised GrabCut: Note that the approach above falls into the category of semi-supervised (interactive) approaches, since it requires bounding boxes as input. Also it allows for binary segmentation only. We will make two simple extensions:

1. First, we want to provide the initial segmentation to the GrabCut algorithm based on an unsupervised initialization. For many input images, a simple clustering algorithm will suffice as initialization.
2. Second, we want to segment an image into more than just two segments. To that end, we can apply α -expansion to the binary problem, in order to find solutions with multiple labels.

¹C. Rother, V. Kolmogorov, and A. Blake. “GrabCut”: Interactive foreground extraction using iterated graph cuts. ACM Transactions on Graphics (SIGGRAPH), August 2004.




Figure 1: Iterative segmentation into three segments. (a) Input image. (b) Initial clustering of color statistics. (c) Final result of iterative graph cuts refinement.

We have provided you with an initial script `problem1` outlining the necessary steps.

Tasks:

- As the first step, we want to find an adequate initialization for the GrabCut algorithm. To that end, please use an algorithm from the `scipy.cluster.vq` package to find `num_segments` initial clusters by means of the color statistics in the input image:

```
cluster_coarse_color_statistics(im, num_segments)
```

Your output should be an integer label mask of the same spatial dimensions as the input image containing labels in $\{0, \text{num_segments} - 1\}$.

4 points

- Implement a function super-imposing labels on top of a color image, *c.f.* Fig. 1:

```
label2color(im, label)
```

To that end, (1) convert the color image into a different color space allowing for separate treatment of color and texture. (2) set the texture from the given color image. (3) set color channels (*e.g.*, *hue*) by means of the labels. Subsequently, assemble the channels in the alternative color space and convert it to back to RGB to form a color-coded representation. For the color-coding, please discretize an existing color map from `matplotlib.cm` by means of the given labels.

3 points

- We already found the edges for the 4-connected Potts model for you. However, we still need the contrast sensitive weights to implement Eq. (3). Please compute the weights in:

```
contrast_weight(im, edges)
```

Set the missing parameter β according to

$$\beta = \left(\frac{1}{|\text{Ed}|} \sum_{(i,j) \in \text{Ed}} \|I_i - I_j\|_2^2 \right)^{-1},$$

where $|\text{Ed}|$ is the number of edges in the Potts model.

3 points

We will now implement some helper functions:

- For the underlying graph cut algorithm, implement the function

```
make_pairwise(lmbda, edges, cweights, num_sites)
```

which constructs a pairwise capacity matrix from the contrast-sensitive Potts model. Note that this step is similar to the last assignment, however, we apply additional weights extracted from color differences.

2 points

- For the unary terms we need to fit and evaluate Gaussian mixture models for the color statistics of the distinct segments. To that end, implement

```
negative_logprob_gmm(im, labels, gmm_components, num_segments)
```

which fits a Gaussian mixture model to the pixels assigned to a segment, and subsequently evaluates the corresponding negative log likelihoods of *all* pixels. Here, `gmm_components` is the number of components to use for the mixture model. Note that you need to fit a Gaussian mixture model for every segment, respectively.

3 points

We now have the ingredients to implement our multi-label segmentation algorithm. Please implement the following two functions, representing the *inner* and *outer* part of the energy minimization:

- Implement a single expansion move, *i.e.* the inner binary update of the algorithm in

```
expand_alpha(alpha, im, label, pairwise, gmm_components, num_segments),
```

where `alpha` is the current label to be expanded, `im` is the input image, `label` is the current assignment of labels, and `pairwise` is the pairwise capacity matrix. Here, you should call `negative_logprob_gmm` to fit and evaluate Gaussian mixture models for all segments, and subsequently construct appropriate unaries from the evaluated (negative) log probabilities. Based on the unaries and pairwise capacities you should update the `label` mask by expanding `alpha` via the graph cuts optimizer provided in `gco.py`.

5 points

- Implement the outer update of the alpha expansion algorithm in

```
iterated_graphcuts(im, label0, pairwise, gmm_components, num_segments)
```

where `label0` is the initial labeling obtained from unsupervised clustering. Here you should repeatedly call `expand_alpha` to perform updates of the current label map. Please find an appropriate configuration (iterations and order of labels) to yield a satisfying labeling.

4 points

- Find an appropriate setting of `gmm_components` and `lmbda` to achieve good results for the given elk image *c.f.* Fig. 1. Please write intermediate labelings of your iterations into a binary folder `bin`. When you submit the assignment, please include your initial labeling `init.png`, an intermediate labeling `intermediate.png`, and your final estimate `final.png`. Do not upload your whole `binary` folder!

3 points