

Computer Vision II

Numpy/PyTorch Introduction

05.06.2019

Why Numpy?

- Fundamental package for scientific computing with Python
 - Powerful **N-dimensional array** objects (`ndarray`)
 - Sophisticated broadcasting (similar, but more intuitive than Julia)
 - Tools for integrating C/C++ code
 - Linear algebra, Fourier transforms, random number generation
- Shared by Python deep learning framework to **communicate with CPU hosts**
 - Cuda/GPU computations → TensorFlow/PyTorch/MXNet datatypes
 - Transfer between GPU and CPU → `Numpy.ndarray`



```
import numpy as np
```

Creating Vectors

MATLAB

PYTHON

JULIA

Row vector: size (1, n)

A = [1 2 3]

A = np.array([1, 2, 3]).reshape(1, 3)

A = [1 2 3]

Column vector: size (n, 1)

A = [1; 2; 3]

A = np.array([1, 2, 3]).reshape(3, 1)

A = [1 2 3]'

1d array: size (n,)

Not possible

A = np.array([1, 2, 3])

A = [1; 2; 3]

or

A = [1, 2, 3]

Integers from j to n with step size k

A = j:k:n

A = np.arange(j, n+1, k)

A = j:k:n

Linearly spaced vector of k points

A = linspace(1, 5, k)

A = np.linspace(1, 5, k)

A = range(1, 5, length = k)

MATLAB

PYTHON

JULIA

Create a matrix

```
A = [1 2; 3 4]
```

```
A = np.array([[1, 2], [3, 4]])
```

```
A = [1 2; 3 4]
```

2 x 2 matrix of zeros

```
A = zeros(2, 2)
```

```
A = np.zeros((2, 2))
```

```
A = zeros(2, 2)
```

2 x 2 matrix of ones

```
A = ones(2, 2)
```

```
A = np.ones((2, 2))
```

```
A = ones(2, 2)
```

2 x 2 identity matrix

```
A = eye(2, 2)
```

```
A = np.eye(2)
```

```
A = I # will adopt  
# 2x2 dims if demanded by  
# neighboring matrices
```

Diagonal matrix

```
A = diag([1 2 3])
```

```
A = np.diag([1, 2, 3])
```

```
A = Diagonal([1, 2,  
3])
```

Uniform random numbers

```
A = rand(2, 2)
```

```
A = np.random.rand(2, 2)
```

```
A = rand(2, 2)
```

MATLAB

PYTHON

JULIA

Access one element

A(2, 2)

A[1, 1]

A[2, 2]

Access specific rows

A(1:4, :)

A[0:4, :]

A[1:4, :]

Access specific columns

A(:, 1:4)

A[:, 0:4]

A[:, 1:4]

Remove a row

A([1 2 4], :)

A[[0, 1, 3], :]

A[[1, 2, 4], :]

Diagonals of matrix

diag(A)

np.diag(A)

diag(A)

Get dimensions of matrix

[nrow ncol] = size(A)

nrow, ncol = np.shape(A)

nrow, ncol = size(A)

- In Python/Numpy indices start with 0 !!!

MATLAB

PYTHON

JULIA

Transpose

`A.'`

`A.T`

`transpose(A)`

Complex conjugate transpose (Adjoint)

`A'`

`A.conj()`

`A'`

Concatenate horizontally

`A = [[1 2] [1 2]]`

or

`A = horzcat([1 2], [1 2])`

`B = np.array([1, 2])`
`A = np.hstack((B, B))`

`A = [[1 2] [1 2]]`

or

`A = hcat([1 2], [1 2])`

Concatenate vertically

`A = [[1 2]; [1 2]]`

or

`A = vertcat([1 2], [1 2])`

`B = np.array([1, 2])`
`A = np.vstack((B, B))`

`A = [[1 2]; [1 2]]`

or

`A = vcat([1 2], [1 2])`

Reshape (to 5 rows, 2 columns)

`A = reshape(1:10, 5, 2)`

`A = A.reshape(5, 2)`

`A = reshape(1:10, 5, 2)`

Convert matrix to vector

`A(:)`

`A = A.flatten()`

`A[:]`

MATLAB

PYTHON

JULIA

Comment one line

% This is a comment

This is a comment

This is a comment

Comment block

```
%{
Comment block
%}
```

```
# Block
# comment
# following PEP8
```

```
#=Comment block
=#
```

For loop

```
for i = 1:N
    % do something
end
```

```
for i in range(n):
    # do something
```

```
for i in 1:N
    # do something
end
```

While loop

```
while i <= N
    % do something
end
```

```
while i <= N:
    # do something
```

```
while i <= N
    # do something
end
```

If

```
if i <= N
    % do something
end
```

```
if i <= N:
    # do something
```

```
if i <= N
    # do something
end
```

If / else

```
if i <= N
    % do something
else
    % do something else
end
```

```
if i <= N:
    # do something
else:
    # so something else
```

```
if i <= N
    # do something
else
    # do something else
end
```

Print text and variable

```
x = 10
fprintf('x = %d \n', x)
```

```
x = 10
print(f'x = {x}')
```

```
x = 10
println("x = $x")
```

Function: anonymous

```
f = @(x) x^2
```

```
f = lambda x: x**2
```

```
f = x -> x^2
# can be rebound
```

Function

```
function out = f(x)
    out = x^2
end
```

```
def f(x):
    return x**2
```

```
function f(x)
    return x^2
end

f(x) = x^2 # not anon!
```

Tuples

```
t = {1 2.0 "test"}
t{1}
```

```
t = (1, 2.0, "test")
t[0]
```

```
t = (1, 2.0, "test")
t[1]
```

Can use cells but watch performance

Closures

```
a = 2.0
f = @(x) a + x
f(1.0)
```

```
a = 2.0
def f(x):
    return a + x
f(1.0)
```

```
a = 2.0
f(x) = a + x
f(1.0)
```

Inplace Modification

```
function f(out, x)
    out = x.^2
end
x = rand(10)
y = zeros(length(x), 1)
f(y, x)
```

```
def f(x):
    x **=2
    return
x = np.random.rand(10)
f(x)
```

```
function f!(out, x)
    out .= x.^2
end
x = rand(10)
y = similar(x)
f!(y, x)
```



Why PyTorch?



Tomasz Malisiewicz
@quantombone

Follow

ICLR 2016 made it seem like everybody was transitioning to TensorFlow, but from my ICLR 2017 convos, PyTorchers are having more fun.

5:40 AM - 25 Apr 2017

58 Retweets 190 Likes



Ferenc Huszár
@fhuszar

Follow

I realised I like @pytorch because it's not a deeplearning framework, it is truly a lightweight, no-brainer numpy replacement and extension

11:12 AM - 4 Oct 2017

47 Retweets 241 Likes



Andrej Karpathy ✅
@karpathy

Following

I've been using PyTorch a few months now and I've never felt better. I have more energy. My skin is clearer. My eye sight has improved.

11:56 AM - 26 May 2017

352 Retweets 1,365 Likes



Goals of deep learning frameworks

- Easily build **big computational graphs**
- Easily **compute gradients** in computational graphs
 - PyTorch utilizes dynamic computation graphs
- Run it all **efficiently on GPU**
 - Wrap thirdparty accelerators such as cuDNN, cuBLAS, etc

PyTorch ~ Python/Numpy

- PyTorch is engineered to be fairly close to Numpy
 - Data types, operation names, arguments etc

Array creation

Numpy	PyTorch	Notes
<code>np.empty((2, 2))</code>	<code>torch.empty(5, 3)</code>	empty array
<code>np.random.rand(3,2)</code>	<code>torch.rand(5, 3)</code>	random
<code>np.zeros((5,3))</code>	<code>torch.zeros(5, 3)</code>	zeros
<code>np.array([5.3, 3])</code>	<code>torch.tensor([5.3, 3])</code>	from list
<code>np.random.randn(*a.shape)</code>	<code>torch.randn_like(a)</code>	
<code>np.arange(16)</code>	<code>torch.range(0,15)</code>	array starting from 0 ending at 15 (inclusive)

Math operations

Numpy	PyTorch	Notes
<code>x+y</code>	<code>x+y</code> <code>y.add_(x)</code> <code>torch.add(x,y)</code>	addition
<code>np.dot(x,y)</code> <code>np.matmul(x,y)</code>	<code>torch.mm(x,y)</code> <code>x.mm(y)</code>	matrix multiplication
<code>x*y</code>	<code>x*y</code>	element-wise multiplication
<code>np.max(x)</code>	<code>torch.max(x)</code>	
<code>np.argmax(x)</code>	<code>torch.argmax(x)</code>	
<code>x**2</code>	<code>x**2</code>	Element-wise powers

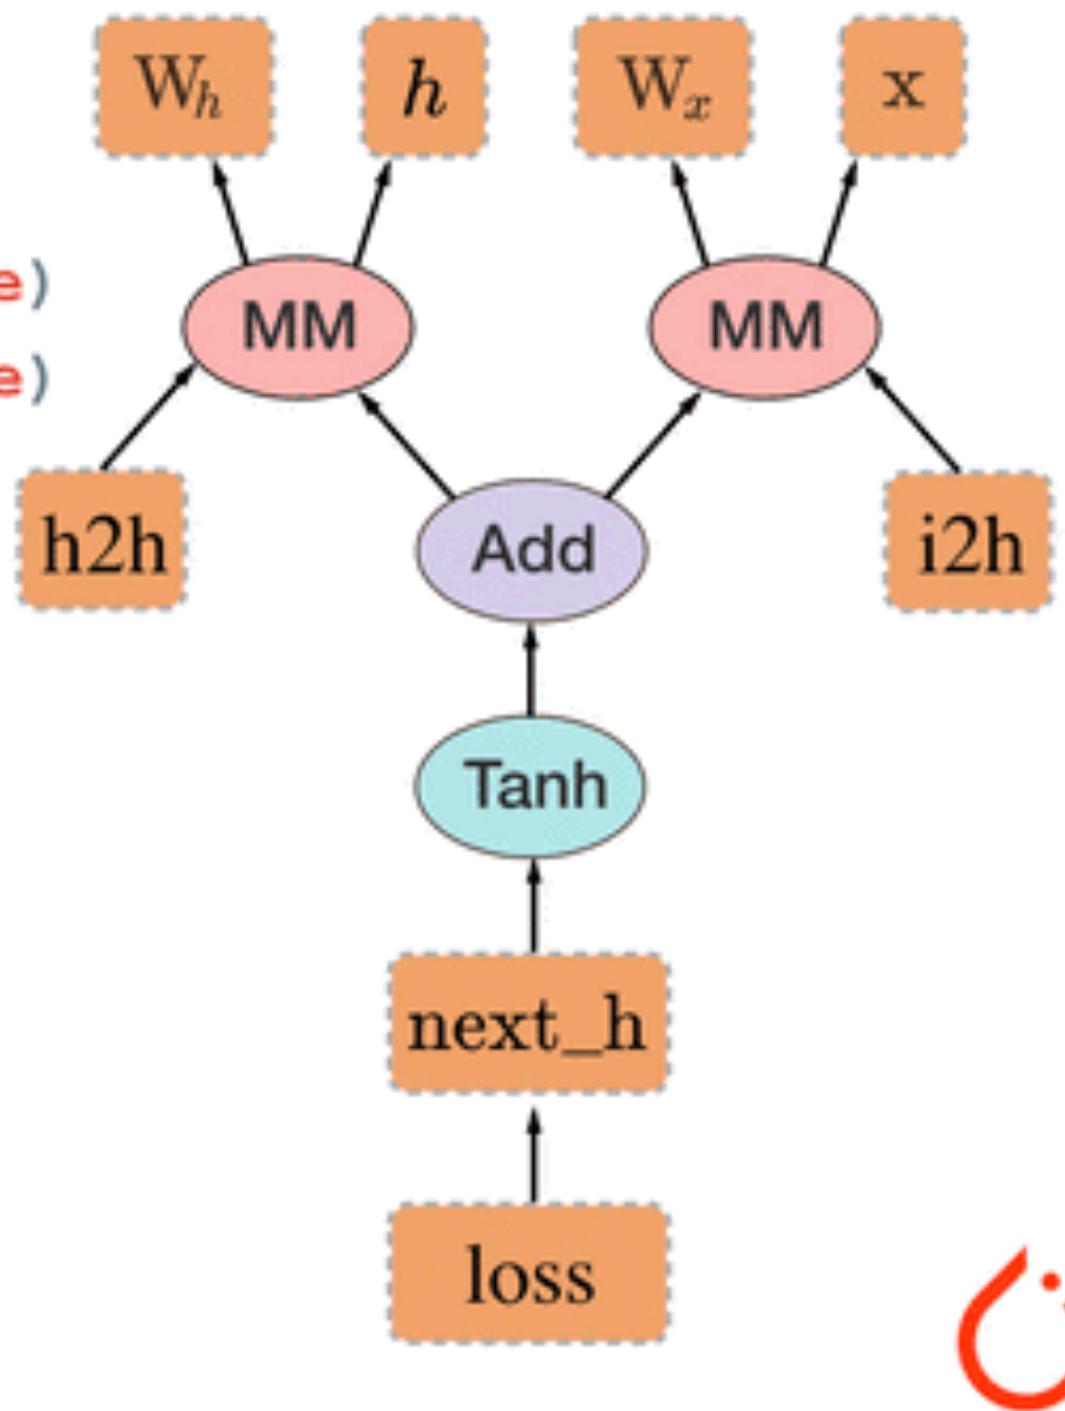
What is a dynamic computation graph?

Back-propagation
uses the dynamically created graph

```
W_h = torch.randn(20, 20, requires_grad=True)
W_x = torch.randn(20, 10, requires_grad=True)
x = torch.randn(1, 10)
prev_h = torch.randn(1, 20)

h2h = torch.mm(W_h, prev_h.t())
i2h = torch.mm(W_x, x.t())
next_h = h2h + i2h
next_h = next_h.tanh()

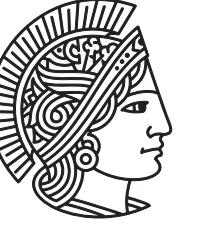
loss = next_h.sum()
loss.backward() # compute gradients!
```





How does this work?

- The key is **automatic differentiation** (Autograd or AD)
- Since we can easily derive derivatives for elementary arithmetic operations – we can also obtain the derivatives for **more complex expressions constructed from elementary operations**
- The mechanics behind this root in the repeated **application of the chain rule**
- Let's look at an example...



Automatic differentiation – An example

- A given scalar function:

$$f = \exp(\exp(x) + \exp(x)^2) + \sin(\exp(x) + \exp(x)^2)$$

- Let's write down the derivative of this:

$$\begin{aligned}\frac{df}{dx} = & \exp(\exp(x) + \exp(x)^2)(\exp(x) + 2\exp(x)^2) \\ & + \cos(\exp(x) + \exp(x)^2)(\exp(x) + 2\exp(x)^2)\end{aligned}$$

Automatic differentiation – An example

- A given scalar function:

$$f = \exp(\exp(x) + \exp(x)^2) + \sin(\exp(x) + \exp(x)^2)$$

- Alternative way is to “program” the expression:

$$a = \exp(x)$$

$$b = a^2$$

$$c = a + b$$

$$d = \exp(c)$$

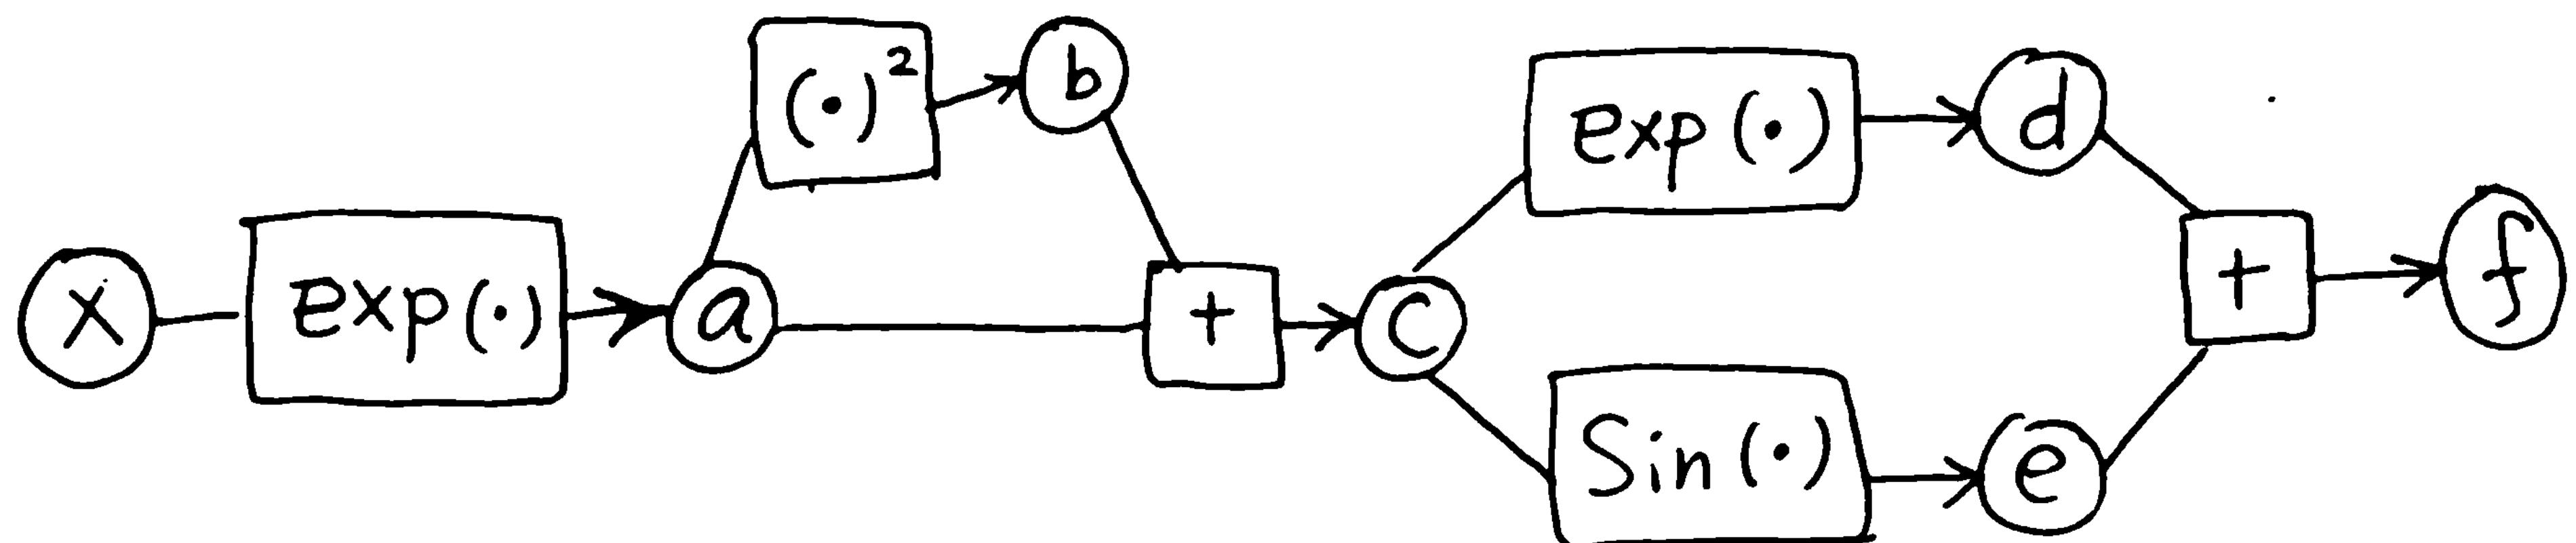
$$e = \sin(c)$$

$$f = d + e$$

Automatic differentiation – An example

$$f = \exp(\exp(x) + \exp(x)^2) + \sin(\exp(x) + \exp(x)^2)$$

- Visualizing the program



$$a = \exp(x)$$

$$b = a^2$$

$$c = a + b$$

$$d = \exp(c)$$

$$e = \sin(c)$$

$$f = d + e$$

Automatic differentiation – An example

- Derive by chaining derivatives of subexpressions

Chain rules

$$\frac{df}{dd} = 1$$

$$\frac{df}{de} = 1$$

$$\frac{df}{dc} = \frac{df}{dd} \frac{dd}{dc} + \frac{df}{de} \frac{de}{dc}$$

$$\frac{df}{db} = \frac{df}{dc} \frac{dc}{db}$$

$$\frac{df}{da} = \frac{df}{dc} \frac{dc}{da} + \frac{df}{db} \frac{db}{da}$$

$$\frac{df}{dx} = \frac{df}{da} \frac{da}{dx}$$



Chain rules

$$\frac{df}{dd} = 1$$

$$\frac{df}{de} = 1$$

$$\frac{df}{dc} = \frac{df}{dd} \exp(c) + \frac{df}{de} \cos(c)$$

$$\frac{df}{db} = \frac{df}{dc}$$

$$\frac{df}{da} = \frac{df}{dc} + \frac{df}{db} 2a$$

$$\frac{df}{dx} = \frac{df}{da} \exp(x).$$



Final expression

$$a = \exp(x)$$

$$b = a^2$$

$$c = a + b$$

$$d = \exp(c)$$

$$e = \sin(c)$$

$$f = d + e$$



Automatic differentiation

- Works for any differential function that can be phrased as an expression graph
- Most important:
 - **Abstracts away difficulties** in gradient computations
 - In turn, this frees you up to **think about the forward computation** and problem at hand
- Pitfalls:
 - Fails if program is not differentiable
 - Difficult to handle if-else blocks
 - Can get rather complicated for vector and matrix expressions
 - e.g. how to compute the derivative of $f = A^{-1}x$?

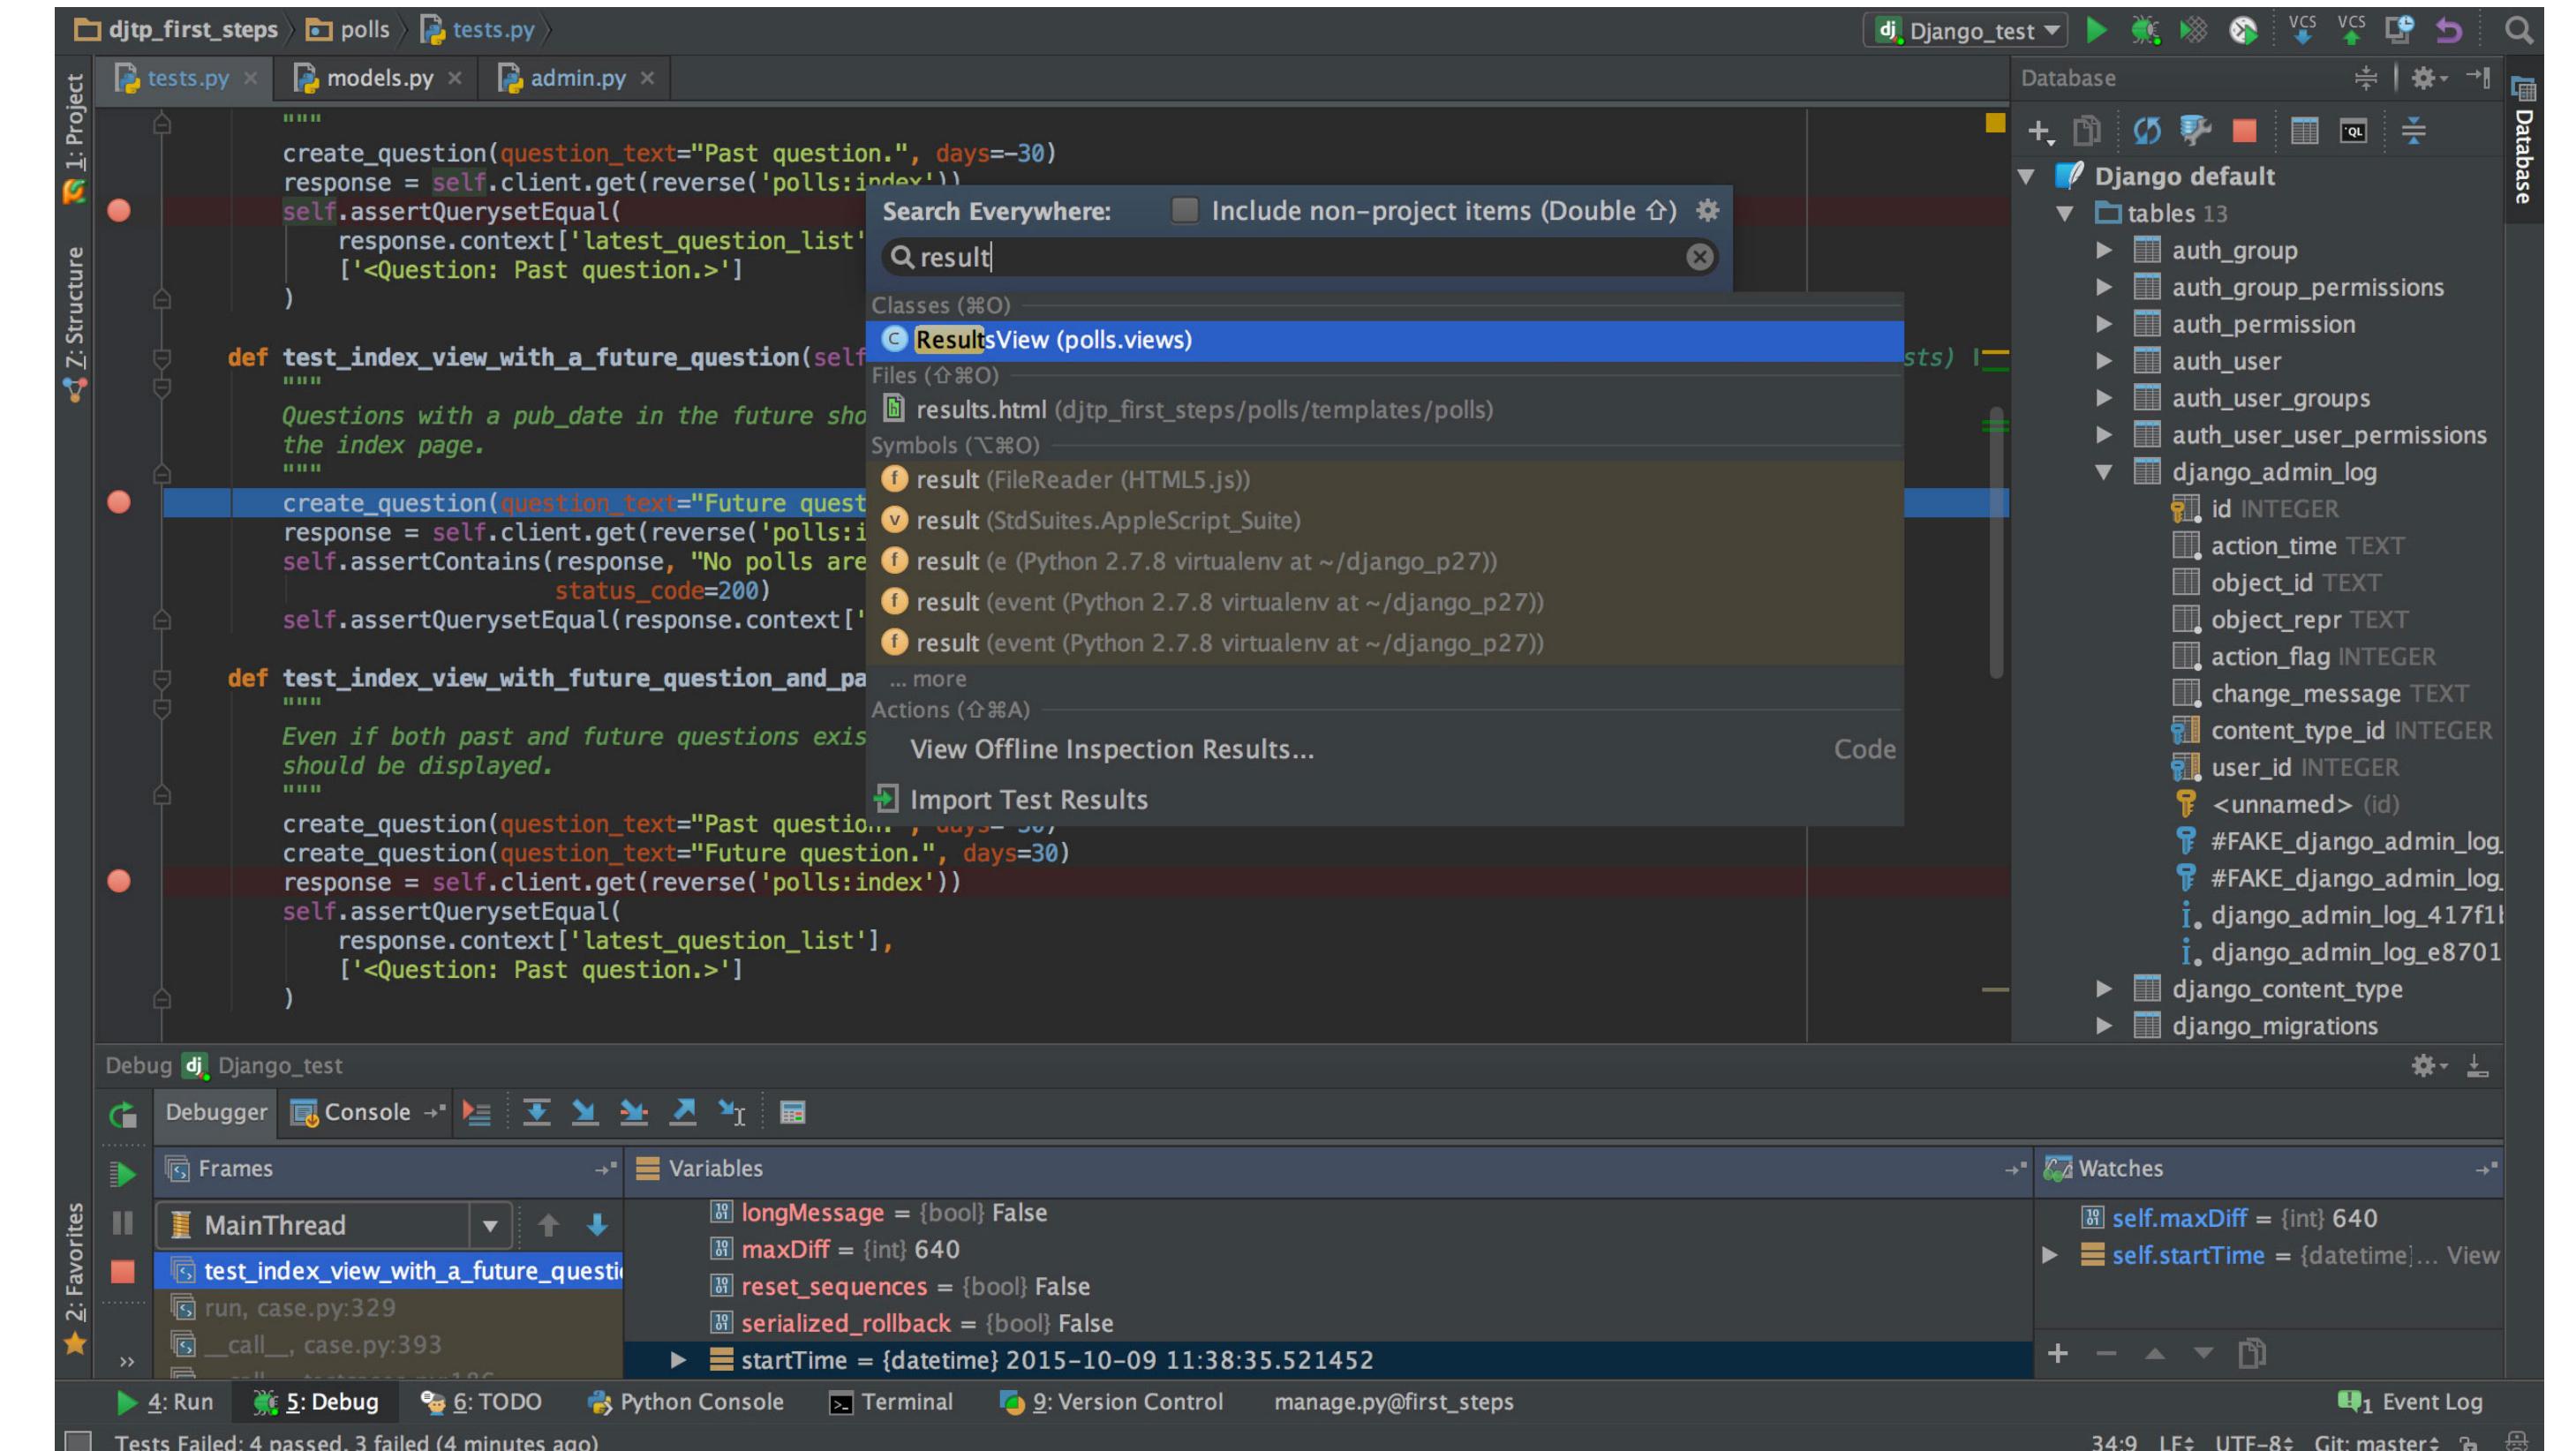
Developing with Python/PyTorch

- Most Popular Python IDE:

 - Code completions
 - Code inspections
 - Visual debugger

- It is usually easiest to tell PyCharm to use an existing Python distribution!

<https://www.jetbrains.com/pycharm/>



How to install Numpy and PyTorch?

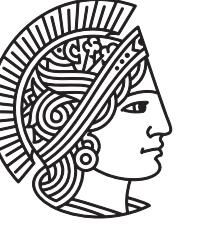
- Install an Anaconda Python3 distribution from:
 - <https://www.anaconda.com/distribution/>
 - Includes comprehensive set of scientific packages (Numpy, Matplotlib, etc)
- Follow instructions from PyTorch webpage:
 - <https://pytorch.org/get-started/locally/>
- This usually just means running:
 - Linux: \$ conda install pytorch-cpu torchvision-cpu -c pytorch
 - Windows: \$ conda install pytorch-cpu torchvision-cpu -c pytorch
 - MacOS: \$ conda install pytorch torchvision -c pytorch

Helpful links

- Good Numpy tutorial:
 - <http://cs231n.github.io/python-numpy-tutorial/>
- PyTorch homepage:
 - <https://pytorch.org/>
- Get started guide
 - <https://pytorch.org/get-started/locally/>
- PyTorch Tutorials
 - <https://pytorch.org/tutorials/>
- PyTorch API:
 - <https://pytorch.org/docs/stable/index.html>

We will now look into the Tutorial notebooks

- Numpy Tutorial:
 - <http://cs231n.github.io/python-numpy-tutorial/>
- What is PyTorch Tutorial
 - https://pytorch.org/tutorials/beginner/blitz/tensor_tutorial.html#sphx-glr-beginner-blitz-tensor-tutorial-py
- Autograd Tutorial:
 - https://pytorch.org/tutorials/beginner/blitz/autograd_tutorial.html#sphx-glr-beginner-blitz-autograd-tutorial-pyet started guide



References

- These slides have been compiled from a variety of resources:
 - <https://cheatsheets.quantecon.org>
 - https://people.cs.umass.edu/~domke/courses/sml/09autodiff_nnets.pdf
 - https://cise.ufl.edu/~xiaoyong/materials/pytorch_tutorial.pdf
 - https://mldublin.github.io/assets/slides/meetup_19/Intro_to_pytorch.html